

Activity: Hands-On-Activity 8.1

Name: Corpuz, Micki Lauren B.

Section: CPE22S3

8.1.1 Intended Learning Outcomes

After this activity, the student should be able to:

- Demonstrate querying and merging of dataframes
- Perform advanced calculations on dataframes
- Aggregate dataframes with pandas and numpy
- Work with time series data

8.1.2 Resources

Computing Environment using Python 3.x

Attached Datasets (under Instructional Materials)

8.1.3 Procedures

The procedures can be found in the canvas module. Check the following under topics:

8.1 Weather Data Collection

8.2 Querying and Merging

8.3 Dataframe Operations

8.4 Aggregations

8.5 Time Series

8.1.4 Data Analysis

Provide some comments here about the results of the procedures

The first procedure shows how to automate the process of collecting daily weather data for one year. The `make_request` function handles the API

communication, while the loop repeatedly calls this function for each day, gathers the data, and stores it in a list.

The second procedure combines two tables. The procedures showed how to query and merge weather data efficiently using pandas. Filtering with `.query()` provided a clean and readable way to isolate specific conditions like snowfall. The merging techniques, including inner, left, right, and outer joins, were used to enrich the weather dataset with station metadata, revealing mismatches and data loss depending on join type. Finally, set operations were helpful in analyzing overlap and differences between datasets before performing merges, making it easier to understand potential data loss or duplication.

The third procedure introduced vast techniques. Z-scores helped identify unusually high trading days for Facebook in 2018, triggered by major news events. Volume change analysis showed market reactions, like the January 12 spike after Facebook's news feed change. Binning trading volume revealed that most days had low activity. For weather data, the `clip()` function kept snowfall days without changing data types. Rolling and expanding windows helped spot trends in weather data. Quantile binning gave better balance than regular binning. Vectorized functions were faster than `iteritems()` for large datasets. Using `pipe()` and custom functions made the code cleaner and more reusable. Lastly, visualization helped illustrate trends in the data.

The fourth procedure introduce methods that show how to summarize, group, and analyze large datasets using pandas and NumPy. For example, we calculated key statistics like total precipitation and snowfall using `.agg()` and `.groupby()` for Central Park and other stations. When analyzing Facebook's stock data, grouping by trading volume helped highlight how price behavior differed across low, medium, and high-volume days. Using the `Groupby` function for aggregating data by time—whether daily, monthly, or quarterly—gave us valuable insights, like identifying which months had the most precipitation. We also used advanced techniques like `transform()`, `pivot_table()`, and `crosstab()` to calculate percentages, reshape data, and count frequencies. These tools are powerful for uncovering patterns and trends in structured datasets.

The last procedure showed techniques to work on with data at different time levels—daily, minute, and second. Using time-based tools, we filtered the data by specific dates and periods. We grouped minute-level data into daily summaries to calculate things like open, high, low, close prices, and volume,

noticing that more trades happen in the first 30 minutes of the day. I used `.shift()` and `.diff()` to measure price changes over time and added new features like after-hours price changes. High-frequency data was resampled to daily and quarterly views to spot broader trends, while missing values were handled with functions like `.pad()` (but now was `.ffill()` --- or forward fill) and `.fillna()`. Lastly, I learned how to merge Facebook and Apple data using time-aware merging methods to align their different timestamp formats accurately.

Also, in this proceure was it where I referenced online as a bunch or warning messages were returned to me about soon-to-be-replaced methods; thus, I tried using the updaed codes here.

8.1.5 Supplementary Activity

Using the CSV files provided and what we have learned so far in this module complete the following exercises:

1. With the earthquakes.csv file, select all the earthquakes in Japan with a magType of mb and a magnitude of 4.9 or greater.

In []:

```
import pandas as pd

data = pd.read_csv('earthquakes.csv')

jp_eq = data[(data['parsed_place']=='Japan') & (data['magType']=='mb') & (data['ma

jp_eq

# Observation:
# The DataFrame jp_eq filters and displays earthquakes in Japan with a magnitude of
```

Out []:

	mag	magType	time	place	tsunami	parsed_place
1563	4.9	mb	1538977532250	293km ESE of Iwo Jima, Japan	0	Japan
2576	5.4	mb	1538697528010	37km E of Tomakomai, Japan	0	Japan
3072	4.9	mb	1538579732490	15km ENE of Hasaki, Japan	0	Japan
3632	4.9	mb	1538450871260	53km ESE of Hitachi, Japan	0	Japan

2. Create bins for each full number of magnitude (for example, the first bin is 0-1, the second is 1-2, and so on) with a magType of ml and count how many are in each bin.

```
In [ ]: df = data.copy()

ml_eq = df[df['magType'] == 'ml']

bins = [b for b in range(11)]

magnitude_counts = pd.cut(ml_eq['mag'], bins=bins, right=False).value_counts().sort
magnitude_counts

# Obsercation
# The counts reflect how frequently earthquakes occur in each magnitude range (Rich
```

Out[]: **count**

mag	
[0, 1)	2072
[1, 2)	3126
[2, 3)	985
[3, 4)	153
[4, 5)	6
[5, 6)	2
[6, 7)	0
[7, 8)	0
[8, 9)	0
[9, 10)	0

dtype: int64

3. Using the faang.csv file, group by the ticker and resample to monthly frequency. Make the following aggregations:

- Mean of the opening price
- Maximum of the high price
- Minimum of the low price
- Mean of the closing price
- Sum of the volume traded

```
In [ ]: df3 = pd.read_csv('faang.csv', index_col='date', parse_dates=True)

monthly_aggregations = df3.groupby('ticker').resample('M').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})

print(monthly_aggregations)

# Observation:
# With ticker on a monthly basis, this computes the average opening and closing price,
# the highest and lowest prices, and the total trading volume for each month.
```

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-31	170.714690	176.6782	161.5708	170.699271	659679440
	2018-02-28	164.562753	177.9059	147.9865	164.921884	927894473
	2018-03-31	172.421381	180.7477	162.4660	171.878919	713727447
	2018-04-30	167.332895	176.2526	158.2207	167.286924	666360147
	2018-05-31	182.635582	187.9311	162.7911	183.207418	620976206
	2018-06-30	186.605843	192.0247	178.7056	186.508652	527624365
	2018-07-31	188.065786	193.7650	181.3655	188.179724	393843881
	2018-08-31	210.460287	227.1001	195.0999	211.477743	700318837
	2018-09-30	220.611742	227.8939	213.6351	220.356353	678972040
	2018-10-31	219.489426	231.6645	204.4963	219.137822	789748068
	2018-11-30	190.828681	220.6405	169.5328	190.246652	961321947
	2018-12-31	164.537405	184.1501	145.9639	163.564732	898917007
AMZN	2018-01-31	1301.377143	1472.5800	1170.5100	1309.010952	96371290
	2018-02-28	1447.112632	1528.7000	1265.9300	1442.363158	137784020
	2018-03-31	1542.160476	1617.5400	1365.2000	1540.367619	130400151
	2018-04-30	1475.841905	1638.1000	1352.8800	1468.220476	129945743
	2018-05-31	1590.474545	1635.0000	1546.0200	1594.903636	71615299
	2018-06-30	1699.088571	1763.1000	1635.0900	1698.823810	85941510
	2018-07-31	1786.305714	1880.0500	1678.0600	1784.649048	97629820
	2018-08-31	1891.957826	2025.5700	1776.0200	1897.851304	96575676
	2018-09-30	1969.239474	2050.5000	1865.0000	1966.077895	94445693
	2018-10-31	1799.630870	2033.1900	1476.3600	1782.058261	183228552
	2018-11-30	1622.323810	1784.0000	1420.0000	1625.483810	139290208
	2018-12-31	1572.922105	1778.3400	1307.0000	1559.443158	154812304
FB	2018-01-31	184.364762	190.6600	175.8000	184.962857	495655736
	2018-02-28	180.721579	195.3200	167.1800	180.269474	516621991
	2018-03-31	173.449524	186.1000	149.0200	173.489524	996232472
	2018-04-30	164.163557	177.1000	150.5100	163.810476	751130388
	2018-05-31	181.910509	192.7200	170.2300	182.930000	401144183
	2018-06-30	194.974067	203.5500	186.4300	195.267619	387265765
	2018-07-31	199.332143	218.6200	166.5600	199.967143	652763259
	2018-08-31	177.598443	188.3000	170.2700	177.491957	549016789
	2018-09-30	164.232895	173.8900	158.8656	164.377368	500468912
	2018-10-31	154.873261	165.8800	139.0300	154.187826	622446235
	2018-11-30	141.762857	154.1300	126.8500	141.635714	518150415
	2018-12-31	137.529474	147.1900	123.0200	137.161053	558786249
GOOG	2018-01-31	1127.200952	1186.8900	1045.2300	1130.770476	28738485
	2018-02-28	1088.629474	1174.0000	992.5600	1088.206842	42384105
	2018-03-31	1096.108095	1177.0500	980.6400	1091.490476	45430049
	2018-04-30	1038.415238	1094.1600	990.3700	1035.696190	41773275
	2018-05-31	1064.021364	1110.7500	1006.2900	1069.275909	31849196
	2018-06-30	1136.396190	1186.2900	1096.0100	1137.626667	32103642
	2018-07-31	1183.464286	1273.8900	1093.8000	1187.590476	31953386
	2018-08-31	1226.156957	1256.5000	1188.2400	1225.671739	28820379
	2018-09-30	1176.878421	1212.9900	1146.9100	1175.808947	28863199
	2018-10-31	1116.082174	1209.9600	995.8300	1110.940435	48496167
	2018-11-30	1054.971429	1095.5700	996.0200	1056.162381	36735570
	2018-12-31	1042.620000	1124.6500	970.1100	1037.420526	40256461
NFLX	2018-01-31	231.269286	286.8100	195.4200	232.908095	238377533
	2018-02-28	270.873158	297.3600	236.1100	271.443684	184585819
	2018-03-31	312.712857	333.9800	275.9000	312.228095	263449491
	2018-04-30	309.129529	338.8200	271.2239	307.466190	262064417
	2018-05-31	329.779759	356.1000	305.7300	331.536818	142051114
	2018-06-30	384.557595	423.2056	352.8200	384.133333	244032001

2018-07-31	380.969090	419.7700	328.0000	381.515238	305487432
2018-08-31	345.409591	376.8085	310.9280	346.257826	213144082
2018-09-30	363.326842	383.2000	335.8300	362.641579	170832156
2018-10-31	340.025348	386.7999	271.2093	335.445652	363589920
2018-11-30	290.643333	332.0499	250.0000	290.344762	257126498
2018-12-31	266.309474	298.7200	231.2300	265.302368	234304628

```
<ipython-input-24-389293fac8f2>:3: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.
```

```
monthly_aggregations = df3.groupby('ticker').resample('M').agg({
```

4. Build a crosstab with the earthquake data between the tsunami column and the magType column. Rather than showing the frequency count, show the maximum magnitude that was observed for each combination. Put the magType along the columns.

```
In [ ]: df4 = data.copy()

crosstab_max_magnitude = pd.crosstab(df4['tsunami'], df4['magType'], values=df4['magType'], aggfunc='max')

crosstab_max_magnitude

# Observation
# The crosstab_max_magnitude table has rows representing whether a tsunami occurred
# This gives a hint that some magTypes may be conducive in tsunami-generating quakes
```

```
Out [ ]: magType  mb  mb_lg  md  mh  ml  ms_20  mw  mwb  mwr  mww
tsunami
0      5.6    3.5  4.11  1.1  4.2   NaN  3.83  5.8  4.8  6.0
1      6.1   NaN  NaN  NaN  5.1   5.7  4.41  NaN  NaN  7.5
```

5. Calculate the rolling 60-day aggregations of OHLC data by ticker for the FAANG data. Use the same aggregations as exercise no. 3.

```
In [ ]: df5 = pd.read_csv('faang.csv', index_col = 'date', parse_dates=True)

rolling_60_day_aggregations = df5.groupby('ticker').rolling(window='60D').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})

rolling_60_day_aggregations

# Observations:
# This shows 60-day rolling summaries of FAANG stock data, where each day's stats are
```

Out[]:

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-02	166.927100	169.0264	166.0442	168.987200	25555934.0
	2018-01-03	168.089600	171.2337	166.0442	168.972500	55073833.0
	2018-01-04	168.480367	171.2337	166.0442	169.229200	77508430.0
	2018-01-05	168.896475	172.0381	166.0442	169.840675	101168448.0
	2018-01-08	169.324680	172.2736	166.0442	170.080040	121736214.0
...
NFLX	2018-12-24	283.509250	332.0499	233.6800	281.931750	525657894.0
	2018-12-26	281.844500	332.0499	231.2300	280.777750	520444588.0
	2018-12-27	281.070488	332.0499	231.2300	280.162805	532679805.0
	2018-12-28	279.916341	332.0499	231.2300	279.461341	521968250.0
	2018-12-31	278.430769	332.0499	231.2300	277.451410	476309676.0

1255 rows × 5 columns

6. Create a pivot table of the FAANG data that compares the stocks. Put the ticker in the rows and show the averages of the OHLC and volume traded data.

```
In [ ]: df6 = pd.read_csv("faang.csv", index_col='date', parse_dates=True)

columns = ['open', 'high', 'low', 'close', 'volume'] # Define the columns for which
pivot_table_faang = df6.pivot_table(index='ticker', values=columns, aggfunc='mean')

pivot_table_faang

# Observation
# This shows the average values for each of the key stock metrics for each ticker.
# Each row represents a different stock like 'AAPL', 'GOOGL', or 'FB', and the columns
# This quickly compares each FAANG stock in terms of price and trading volume over
```


Out[]:

	close	high	low	open	volume
ticker					
AAPL	186.986218	188.906858	185.135729	187.038674	3.402145e+07
AMZN	1641.726175	1662.839801	1619.840398	1644.072669	5.649563e+06
FB	171.510936	173.615298	169.303110	171.454424	2.768798e+07
GOOG	1113.225139	1125.777649	1101.001594	1113.554104	1.742645e+06
NFLX	319.290299	325.224583	313.187273	319.620533	1.147030e+07

7. Calculate the Z-scores for each numeric column of Netflix's data (ticker is NFLX) using apply().

```
In [ ]: from scipy.stats import zscore
```

```
In [ ]: df7 = df6.copy()

nflx = df7[df['ticker']=='NFLX']

numeric_columns = nflx.select_dtypes(include='number')

z_scores = numeric_columns.apply(zscore)

z_scores

# Observation:
# This tells us whether certain stock values are much higher or lower than usual
```

Out[]:

	open	high	low	close	volume
date					
2018-01-02	-2.505749	-2.521050	-2.415042	-2.421473	-0.088937
2018-01-03	-2.385047	-2.428022	-2.290360	-2.339951	-0.508620
2018-01-04	-2.300860	-2.410885	-2.239081	-2.328071	-0.961204
2018-01-05	-2.279559	-2.350294	-2.206487	-2.238767	-0.783894
2018-01-08	-2.223367	-2.299699	-2.148042	-2.196572	-1.040606
...
2018-12-24	-1.574618	-1.521399	-1.630448	-1.749435	-0.339680
2018-12-26	-1.738529	-1.442855	-1.680690	-1.344082	0.518073
2018-12-27	-1.410097	-1.420618	-1.498794	-1.305267	0.135138
2018-12-28	-1.251257	-1.291594	-1.299877	-1.294718	-0.085334
2018-12-31	-1.206222	-1.124597	-1.090706	-1.057529	0.360163

251 rows × 5 columns

8. Add event descriptions:

- Create a dataframe with the following three columns: ticker, date, and event. The columns should have the following values:
- ticker: 'FB'
- date: ['2018-07-25', '2018-03-19', '2018-03-20']
- event: ['Disappointing user growth announced after close.', 'Cambridge Analytica story', 'FTC investigation']
- Set the index to ['date', 'ticker']
- Merge this data with the FAANG data using an outer join

```
In [ ]: event_data = {
    'ticker': ['FB', 'FB', 'FB'],
    'date': ['2018-07-25', '2018-03-19', '2018-03-20'],
    'event': ['Disappointing user growth announced after close.', 'Cambridge Analyt
}
events_df = pd.DataFrame(event_data)

events_df['date'] = pd.to_datetime(events_df['date']) # Convert date column to date
events_df.drop_duplicates(inplace=True) # Remove duplicates
events_df.set_index(['date', 'ticker'], inplace=True)

faang_df = pd.read_csv("faang.csv", parse_dates=['date'])# Read the FAANG data into
```

```
faang_df.set_index(['date', 'ticker'], inplace=True)

merged_df = faang_df.merge(events_df, how='left', left_index=True, right_index=True)


merged_df

# Observation:
# This tells us how events like the Cambridge Analytica scandal or poor user growth
# and other financial metrics to identify correlations or trends.
```

Out[]:

		open	high	low	close	volume	event
date	ticker						
2018-01-02	FB	177.68	181.58	177.5500	181.42	18151903	NaN
2018-01-03	FB	181.88	184.78	181.3300	184.67	16886563	NaN
2018-01-04	FB	184.90	186.21	184.0996	184.33	13880896	NaN
2018-01-05	FB	185.59	186.90	184.9300	186.85	13574535	NaN
2018-01-08	FB	187.20	188.90	186.3300	188.28	17994726	NaN
...
2018-12-24	GOOG	973.90	1003.54	970.1100	976.22	1590328	NaN
2018-12-26	GOOG	989.01	1040.00	983.0000	1039.46	2373270	NaN
2018-12-27	GOOG	1017.15	1043.89	997.0000	1043.88	2109777	NaN
2018-12-28	GOOG	1049.62	1055.56	1033.1000	1037.08	1413772	NaN
2018-12-31	GOOG	1050.96	1052.70	1023.5900	1035.61	1493722	NaN

1255 rows × 6 columns

 Trial only

```
In [ ]: data1 = merged_df[merged_df['event'] == 'Disappointing user growth announced after']
data2 = merged_df[merged_df['event'] == 'Cambridge Analytica story']
data3 = merged_df[merged_df['event'] == 'FTC investigation']
```

```
In [ ]: # Concatenate data1, data2, and data3 into one DataFrame
merged_data = pd.concat([data1, data2, data3])

merged_data
```

Out[]:

		open	high	low	close	volume	event
date	ticker						
2018-07-25	FB	215.715	218.62	214.27	217.50	64592585	Disappointing user growth announced after close.
2018-03-19	FB	177.010	177.17	170.06	172.56	88140060	Cambridge Analytica story
2018-03-20	FB	167.470	170.20	161.95	168.15	129851768	FTC investigation

9. Use the `transform()` method on the FAANG data to represent all the values in terms of the first date in the data. To do so, divide all the values for each ticker by the values

for the first date in the data for that ticker. This is referred to as an index, and the data for the first date is the base (<https://ec.europa.eu/eurostat/statistics-explained/index.php/Beginners:Statisticalconcept-Indexandbaseyear>). When data is in this format, we can easily see growth over time. Hint: `transform()` can take a function name.

```
In [ ]: faang_df = pd.read_csv("faang.csv", parse_dates=['date'])

# Select only numeric columns
numeric_cols = faang_df.select_dtypes(include='number').columns

# Define the indexing function
def first_date_index(group):
    return group / group.iloc[0]

faang_indexed_df = faang_df.copy()
faang_indexed_df[numeric_cols] = faang_df.groupby('ticker')[numeric_cols].transform(
    first_date_index)

faang_indexed_df

# Observation:
# This normalizes the values for each FAANG stock, where all numeric columns are di
# That is why everything starts at 1.0. This makes it easy to compare growth or dec
# This tells how each stock has performed relative to its own starting point.
```

Out[]:

	ticker	date	open	high	low	close	volume
0	FB	2018-01-02	1.000000	1.000000	1.000000	1.000000	1.000000
1	FB	2018-01-03	1.023638	1.017623	1.021290	1.017914	0.930292
2	FB	2018-01-04	1.040635	1.025498	1.036889	1.016040	0.764707
3	FB	2018-01-05	1.044518	1.029298	1.041566	1.029931	0.747830
4	FB	2018-01-08	1.053579	1.040313	1.049451	1.037813	0.991341
...
1250	GOOG	2018-12-24	0.928993	0.940578	0.928131	0.916638	1.285047
1251	GOOG	2018-12-26	0.943406	0.974750	0.940463	0.976019	1.917695
1252	GOOG	2018-12-27	0.970248	0.978396	0.953857	0.980169	1.704782
1253	GOOG	2018-12-28	1.001221	0.989334	0.988395	0.973784	1.142383
1254	GOOG	2018-12-31	1.002499	0.986653	0.979296	0.972404	1.206986

1255 rows × 7 columns

Conclusion

In this activity, I worked with different ways to aggregate data using Pandas. I performed actions like `groupby()`, `resample()`, `pivot_table()`, and `crosstab()` to calculate averages, totals, and the distribution of data. I also used techniques like rolling windows and z-score normalization to make the data easier to compare, and I merged event data with stock data to explore how certain events might have affected stock prices.

I also tried looking online for references, especially since I was trying out some new Pandas functions (for me so to speak). Plus, I wanted to double-check if they were still being used the same way in the current version, which I found by looking at some GitHub examples.

Overall, this activity helped me get more techniques with analyzing and summarizing data.